

# Scalable Sales GenAI Architecture

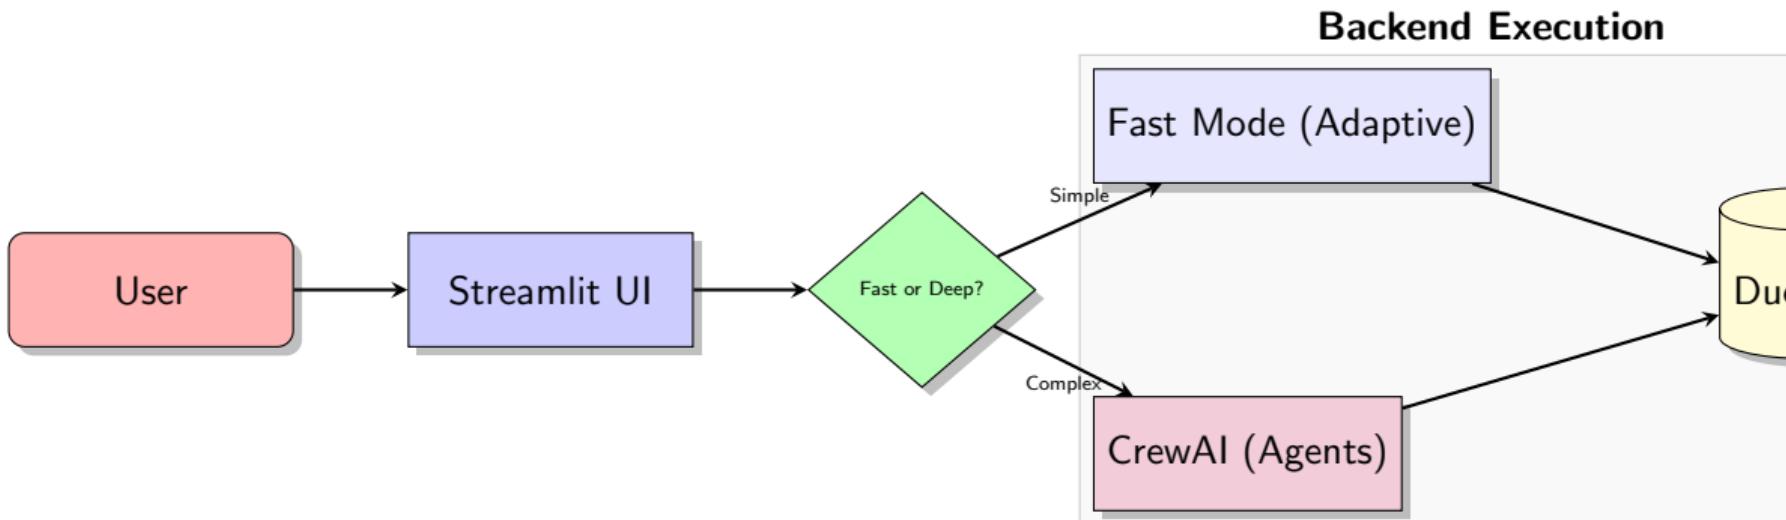
## Production-Ready CrewAI + OpenAI + DuckDB

Abhishek Dangwal

February 8, 2026

*Architecture Overview for 100GB Scale Implementation*

# High-Level System Architecture



## Key Components:

- **UI:** Streamlit with structured rendering (SQL + Data + Insights).
- **Orchestration:** CrewAI for agent delegation; Fast path for low latency.
- **Data:** Persistent DuckDB using .cache/ storage.

# LLM Integration Strategy

## Fast Mode (Low Latency)

- **Goal:** Speed & Low Cost.
- **Technique:** Adaptive Schema-Driven SQL Synthesis.
- **Model:** gpt-4o-mini.
- **Use Case:** Simple lookups, standard aggregations.
- *Bypasses heavy agent loops.*

## Deep Mode (High Fidelity)

- **Goal:** Complex Reasoning.
- **Technique:** CrewAI Multi-Agent System.
- **Workflow:**
  - Schema Chunking (Reduce context).
  - SQL Validation Agent.
  - Auto-Repair loop.

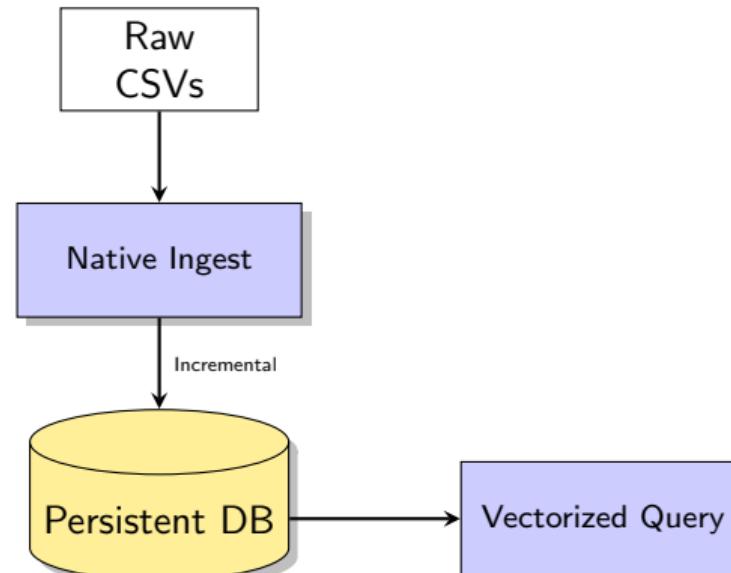
# Data Storage & Retrieval: 100GB Scale Design

## DuckDB Architecture Strategy:

- **Persistent Storage:** Data stored in `.cache/sales.duckdb` to prevent full reloads.
- **Native Ingestion:** Uses `read_csv_auto` (Zero-copy) instead of Pandas loading for lower memory footprint.
- **Dynamic Marts:** Auto-generation of `mart_dyn_*` views for new CSV files.

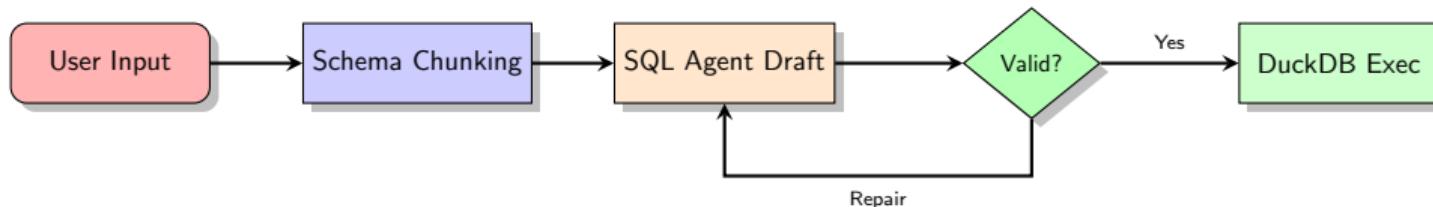
## Performance at Scale:

- Vectorized Execution Engine (OLAP optimized).
- Incremental refresh checks (Time-stamp based).



# Example Query-Response Pipeline

Scenario: "Show top 5 regions by sales in Q3."



## Pipeline Steps

- ① **Schema Chunking:** Retrieve only relevant table definitions (Cost saving).
- ② **Drafting:** Agent generates SQL grounded in Dataset Playbook.
- ③ **Guardrails:** Validator checks for SELECT-only safety and column existence.
- ④ **Visuals:** Streamlit renders SQL code block + DataFrame results.

# Cost and Performance Considerations

Optimization Area	Implementation Strategy
<b>Token Cost</b>	Schema Chunking (Only send relevant headers) Use of gpt-4o-mini for 80% of queries
<b>Latency</b>	<b>Response Caching</b> for repeated queries Fast Mode for non-reasoning tasks
<b>Ingestion Speed</b>	read_csv_auto (DuckDB) vs Pandas Incremental file loading (Timestamp checks)

## Future Improvements:

- Support for Parquet (Compression for >100GB).
- Async background workers for request queuing.

# Thank You

Questions? email me at:[dangwalabhishek5@gmail.com](mailto:dangwalabhishek5@gmail.com)